

<div> <div>Shardul Chavan</div> <div>AWS Certified - Data Engineer Associate</div> <div>+18573135138 shardulc36@gmail.com LinkedIn GitHub</div> </div>	
EDUCATION	
Northeastern University - Master of Science, Information Systems, GPA: 3.76/4.00	Sep 2022 — Dec 2024
University of Mumbai - Bachelor of Engineering, Computer Science and Engineering	Aug 2019 — May 2022
WORK EXPERIENCE	
Crewasis, Remote - Data Scientist Intern	Jan 2025 — Jun 2025
<ul style="list-style-type: none"> ETL Workflow Automation: Designed and deployed scalable ETL pipelines processing 2M+ unstructured text records from social/news feeds, enabling downstream NLP, sentiment analysis, and actionable product-market insights. Cloud ML Orchestration: Built EC2-hosted machine learning pipeline orchestrated with AWS Lambda, automating Latent Semantic Indexing (LSI) training and reducing manual retraining cycles by 70%. Infrastructure-as-Code (IaC): Utilized Terraform scripts to deploy repeatable AWS infrastructure, ensuring high availability and reproducibility. Interactive Reporting & Analytics: Developed interactive Tableau dashboards combining NLP insights with CRM segmentation, empowering marketing teams to make data-driven decisions and increasing campaign conversions by 15%. 	
Skyworks Solutions Inc., Boston, MA - Data Analytics Engineer Co-op	Jan 2024 — Jun 2024
<ul style="list-style-type: none"> Airflow Pipeline Development: Engineered 25+ Airflow and dbt pipelines across Azure SQL and data lakes, reducing data refresh latency by 60% and boosting accessibility for analytics teams. Data Product Optimization: Standardized RF testing data into analytics-ready formats, saving 5+ engineering hours weekly and accelerating high-volume RF module validation. System Resilience Engineering: Containerized data services with Docker and set up Prometheus monitoring to track performance and failures, reducing downtime by 45% through retry logic and better visibility. Ensured data governance by integrating validation with Great Expectations and enforcing RBAC + row-level security policies, ensuring compliance and secure data access. Data Modeling & Reporting: Designed scalable dimensional data models and built interactive Power BI dashboards for 3 business units, accelerating monthly reporting cycles and enabling cross-team insights. 	
Northeastern University, Boston, MA - BI Analyst & AI Researcher	Sep 2023 — Dec 2024
<ul style="list-style-type: none"> Big Data Pipeline Engineering: Developed scalable PySpark pipelines on Hadoop to process 1.2M+ insurance records, enabling efficient feature engineering for risk modeling and improving predictive accuracy by 20%. AI Solution Development: Mentored 50+ graduate students in building applied AI/BI applications, integrating LLMs (OpenAI, Anthropic) into enterprise analytics use cases and accelerating delivery of intelligent insights. Product-facing BI automation: Built SnowGPT a LangChain-powered tool on Snowflake, enabling natural language-to-SQL translation driving 40% increase in self-service analytics adoption among non-technical stakeholders. CI/CD for ML Workflows: Deployed containerized ML models with Docker and GitHub Actions CI/CD, reducing manual deployment time by 60% and ensuring reproducibility across staging and production 	
Accion Labs, Mumbai, India - Data Analytics Engineer	Jan 2022 — Jul 2022
<ul style="list-style-type: none"> LLM Integration & Analytics Automation: Engineered REST-based integrations between LLM APIs and ServiceNow, enhancing conversational AI for support bots, resulting in a 40% reduction in ticket escalations and improved NLU accuracy. SQL Automation & Cross-Team Collaboration: Built parameterized SQL automation scripts to streamline reporting cycles, while partnering with QA and product teams to validate data accuracy against KPIs and SLA thresholds. 	
PROJECTS	
NewsSphere: Personalized News Digest Platform Google Cloud, Docker, Pinecone, AzureSQL, FastAPI	
<ul style="list-style-type: none"> Engineered FastAPI-based services with optimized RAG indexing, reducing content retrieval times by 35% and driving higher user engagement. Deployed containerized Streamlit app and services with Docker + Kubernetes, ensuring fast, reliable personalized news delivery every 8 minutes 	
AWS-Based YouTube Data Analytics IAM, Redshift, S3, Lambda, Spark, Glue, Athena	
<ul style="list-style-type: none"> Built a serverless ELT pipeline using AWS Lambda, Glue, Redshift to ingest multi-regional YouTube data, and analyze 80K+ user interactions in real time, powering QuickSight dashboards that informed product engagement and growth strategies for marketing teams Automated deployment with Terraform IaC and scaled workloads via Kubernetes (EKS), improving pipeline resilience and cost efficiency. 	
Retail Sales Data Pipeline Azure Blob Storage, Databricks, Spark, Azure SQL	
<ul style="list-style-type: none"> Architected a Spark pipeline processing 1M+ retail transactions, enabling trend analysis and accurate forecasting. Implemented Grafana dashboards and OpenTelemetry tracing to proactively detect errors and ensure pipeline observability. 	
Iowa Retail Sales Intelligence & Reporting Alteryx, Talend, Azure Data Factory, Power BI, Tableau, Azure Data Studio	
<ul style="list-style-type: none"> Engineered ETL workflows using Azure Data Factory and Alteryx to process 25M+ sales records, built Kimball-based models to boost ingestion speed by 60%, and developed Power BI dashboards to improve sales forecasting and inventory planning. 	
SKILLS	
Programming & Scripting	: Python, SQL, Scala, Bash, Java, C++
Data Engineering & Orchestration	: Airflow, dbt, Spark (PySpark), Databricks, AWS Glue, Azure Data Factory
Cloud & Infrastructure	: AWS (S3, EC2, Lambda, Redshift, Glue, EKS), Azure (ADF, Synapse), GCP (BigQuery, Cloud Run), Kubernetes, Helm, ArgoCD, Flux, Terraform
Data Modeling & Governance	: Snowflake, Redshift, PostgreSQL, MongoDB, Dimensional Modeling, RBAC, Row-Level Security, Metadata Lineage
Observability & Monitoring	: Prometheus, CloudWatch, OpenTelemetry, Splunk, Grafana, ELK stack
Business Intelligence	: Tableau, Power BI, QuickSight, Streamlit, Superset
APIs & Frameworks	: FastAPI, REST APIs, Flask
DevOps & CI/CD	: Docker, GitHub Actions, Jenkins, CircleCI
Analytics & Experimentation	: A/B Testing, Hypothesis Testing, Product KPIs, Event Logging